

Contribution Title

First Author¹[0000–1111–2222–3333], Second Author^{2,3}[1111–2222–3333–4444], and
Third Author³[2222–3333–4444–5555]

¹ Princeton University, Princeton NJ 08544, USA

² Springer Heidelberg, Tiergartenstr. 17, 69121 Heidelberg, Germany
lncs@springer.com

<http://www.springer.com/gp/computer-science/lncs>

³ ABC Institute, Rupert-Karls-University Heidelberg, Heidelberg, Germany
{abc,lncs}@uni-heidelberg.de

Abstract. The abstract should briefly summarize the contents of the paper in 150–250 words.

Keywords: First keyword · Second keyword · Another keyword.

1 Introduction

Motion planning requires a robot to reach a goal while remaining collision free and respecting its motion constraints. This task is challenging because obstacles make the free configuration space non-convex. Sampling-based planners, such as probabilistic roadmaps and rapidly exploring random trees, can explore complex spaces effectively [2,3]; trajectory-optimization methods can generate smooth motions but commonly depend on a feasible initialization and may converge to a local solution [9,7]. Methods that expose convex structure in the free space offer a complementary route to handling this geometric non-convexity.

Graphs of Convex Sets (GCS) formulate motion planning as a joint discrete and continuous optimization problem [6]. A collision-free free-space cover is represented by convex regions and their feasible transitions; GCS then selects a route through this graph while optimizing the trajectory within the chosen regions. When combined with Bézier curves, the convex-hull property provides continuous-time collision guarantees from convex containment constraints [5]. The quality of the convex cover is therefore central to practical performance. It controls not only the number of regions, but also their overlap pattern, the number of GCS edges, and ultimately the size of the optimization problem.

This paper studies that front-end design choice by comparing Approximate Convex Decomposition (ACD) [4] and Visibility Clique Cover (VCC) [8] in the same GCS–Bézier planning pipeline. ACD uses local concavity to decompose polygonal free space, whereas VCC uses sampled visibility and convex-region inflation to seek a compact cover. These methods embody different trade-offs: ACD can be inexpensive to construct but may yield many local regions, while VCC can reduce the downstream graph but incurs sampling and clique-cover costs and can be sensitive to narrow passages.

The paper makes the following empirical contributions:

- It provides a paired comparison of ACD and VCC on five two-dimensional map families, using five randomized instances per family and a common GCS–Bézier planner.
- It evaluates decomposition time, GCS graph and presolved-model size, solve time, end-to-end time, success rate, and trajectory quality rather than using the number of regions as the sole proxy for performance.
- It shows that the induced edge structure is a stronger predictor of GCS solve time than region count alone. In the tested single-query setting, ACD is the more reliable and faster end-to-end default, while VCC is most useful when its decomposition can be amortized across repeated queries or when it substantially reduces graph connectivity.

The remainder of the paper is organized as follows. Section 2 summarizes the GCS free-space representation and the two decomposition methods. Section 3 presents the GCS–Bézier planning framework. Section 4 reports the experimental evaluation, and Section 5 concludes the paper and outlines future work.

2 Technical Background

2.1 Convex Free-Space Covers for GCS–Bézier

Obstacle avoidance makes motion planning difficult because the free configuration space $\mathcal{C}_{\text{free}}$ is generally non-convex. The GCS–Bézier approach addresses this difficulty by representing the part of $\mathcal{C}_{\text{free}}$ relevant to a query with a finite collection of collision-free convex regions,

$$\mathcal{Q} = \{Q_i\}_{i=1}^{N_R}, \quad Q_i \subseteq \mathcal{C}_{\text{free}}.$$

The regions may overlap. They induce a graph with one vertex per region and an edge for each feasible transition between adjacent regions. The start and goal are connected to the regions that contain them or can be reached from them.

Graphs of Convex Sets (GCS) jointly optimize the discrete route through this graph and the continuous trajectory variables associated with the selected regions [6,5]. In the GCS–Bézier formulation, a trajectory segment is represented by Bézier control points constrained to its assigned region. By the convex-hull property of Bézier curves, this containment certifies collision freedom for the whole segment rather than only for sampled points.

Consequently, decomposition is not merely a preprocessing step. It determines the number of graph vertices N_R , the number of feasible transitions N_E , and the size of the resulting convex optimization model. A cover with fewer regions is not necessarily easier to solve: large or highly overlapping regions can induce a dense transition graph. This motivates evaluating decomposition methods through their downstream GCS model and trajectory, not through region count alone.

2.2 Approximate Convex Decomposition

Approximate Convex Decomposition (ACD) is a geometry-driven method for decomposing a non-convex polygon [4]. It recursively locates the most significant concavity, inserts a cut to resolve it, and repeats until the residual concavity of every component is below a user-selected tolerance τ . The tolerance controls the granularity of the decomposition: a smaller τ resolves more geometric detail and usually produces more regions, whereas a larger τ avoids cuts associated with minor features.

For the present study, ACD provides a direct and inexpensive way to construct the safe-region representation for a planar environment. Its principal trade-off is that resolving local geometric features can produce many regions and, therefore, many possible GCS transitions. This makes ACD a useful baseline for testing whether a faster decomposition stage offsets the cost of a larger downstream graph.

2.3 Visibility Clique Cover

Visibility Clique Cover (VCC) constructs a convex cover from the global visibility structure of the free space [8]. It first samples collision-free configurations and connects pairs whose straight-line segment remains collision free. Cliques in this visibility graph identify groups of mutually visible samples. Each clique initializes an ellipsoid, which is inflated into a collision-free convex polytope using an IRIS-style region-inflation step [1]. Sampling and inflation are repeated until the desired coverage is reached.

Because it groups mutually visible samples before inflation, VCC can produce a small number of large regions and substantially reduce the GCS graph for some environments. Its cost, however, depends on the sampling budget, visibility tests, clique-cover computation, and coverage target. Sparse or unrepresentative samples can also make narrow passages difficult to cover or connect. The comparison in this paper therefore examines the complementary trade-off between ACD’s inexpensive local decomposition and VCC’s potentially more compact but costlier visibility-based cover.

3 Integrated Framework GCS-Bézier for Motion Planning

4 Experimental Results and Discussion

4.1 Benchmark Overview

The benchmark evaluates two automatic free-space decomposition methods, Approximate Convex Decomposition (ACD) and Visibility Clique Cover (VCC), inside the same GCS-Bézier motion-planning pipeline. Five map families are used: `clustering`, `narrows`, `rooms`, `u_shaped`, and `flappy`. Each family is generated

with five seeds under the same workspace scale, obstacle density, obstacle size, and start-goal configuration, giving 50 paired runs in total.

ACD solved all 25 instances. VCC solved 23 out of 25 instances; both failures occurred on the **narrows** family, at seeds 2 and 3. The failed VCC diagnostics still reported moderate graph sizes before termination: 11 vertices and 26 edges for seed 2, and 15 vertices and 57 edges for seed 3. Thus, failure is not explained only by graph size. The more likely cause is the interaction between visibility sampling, coverage requirements, start-goal connectivity, and the narrow-gateway geometry.

4.2 Model Size and Runtime

Table 1 summarizes the main model-size and timing metrics. Each entry is the median over successful seeds. For VCC on **narrows**, the statistics are computed over the three successful runs, while the success rate remains reported as 3/5.

Table 1. Median model size and runtime by map family. Time is measured in seconds.

Map	Method	Success	N_R	N_E	t_{dec}	t_{GCS}	$t_{\text{dec}} + t_{\text{GCS}}$
Clustering	ACD	5/5	104	462	7.73	40.10	59.30
	VCC	5/5	22	99	85.80	3.10	88.80
Narrows	ACD	5/5	13	40	0.68	0.83	1.52
	VCC	3/5	11	50	13.10	0.73	13.80
Rooms	ACD	5/5	15	38	0.78	0.50	1.28
	VCC	5/5	13	57	15.30	1.33	18.90
U-shaped	ACD	5/5	25	86	1.29	2.77	4.07
	VCC	5/5	16	89	16.30	1.93	18.30
Flappy	ACD	5/5	13	26	0.98	0.16	1.13
	VCC	5/5	13	28	12.80	0.16	12.90

VCC is most effective on the **clustering** maps. It reduces the median number of regions from 104 to 22, the number of GCS edges from 462 to 99, and the GCS wall time from 40.10 s to 3.10 s. This confirms that visibility-based grouping can substantially reduce the downstream optimization burden when the free space contains many locally redundant convex pieces.

The same conclusion does not hold uniformly across all map families. On **u_shaped**, VCC reduces the number of regions from 25 to 16, but the median edge count is essentially unchanged. On **rooms**, VCC has an outlier with 33 regions, 430 edges, and a GCS wall time of 105 s. These cases show that the edge structure of the GCS graph matters as much as the number of convex regions.

When decomposition time is included, ACD is faster in all five map families under the tested configuration. The VCC setup uses 500 samples and a target coverage of 0.96, making its decomposition cost much larger than that of ACD. Even for **clustering**, where VCC reduces GCS solve time by roughly one order of magnitude, its median decomposition time of 85.80 s makes the end-to-end runtime higher than ACD. VCC is therefore more attractive when the same

decomposition can be reused for many planning queries; ACD is preferable for one-shot planning in this benchmark.

4.3 Correlation Between Model Size and Solve Time

Figure 1 shows the relationship between model-size indicators and solve time over the 48 successful runs. The strongest Spearman correlations with **SolvePath** time are obtained by the presolved model size: $\rho = 0.986$ for the number of presolved variables and $\rho = 0.983$ for the number of presolved constraints. Among graph-level quantities available before solver construction, the number of edges is more predictive than the number of vertices, with $\rho = 0.982$ compared with $\rho = 0.839$.




Fig. 1. Relationships between model-size indicators and solve time over 48 successful runs. Log axes are used where the metric range spans several orders of magnitude.

This result is consistent with the GCS formulation. Edges do not merely encode topological adjacency; they also introduce flow variables and continuity constraints between adjacent convex sets. A decomposition with fewer regions can still create a difficult optimization problem if it induces a dense graph. Consequently, the number of regions alone is an incomplete proxy for computational difficulty.

4.4 Trajectory Quality

Table 2 reports median path quality metrics. Neither method dominates in path length or trajectory time. ACD produces shorter paths on **clustering** and

u_shaped; VCC performs better on **rooms**, is nearly equivalent on **flappy**, and has only three successful **narrows** runs for comparison.

Table 2. Median trajectory-quality metrics over successful runs.

Map	Method	Path length	Traj. time	Min. clearance	Active regions
Clustering	ACD	7.09	8.82	6.38×10^{-5}	13
	VCC	8.29	8.38	1.66×10^{-1}	5
Narrows	ACD	12.50	15.70	4.79×10^{-4}	9
	VCC	12.00	16.10	2.50×10^{-3}	7
Rooms	ACD	8.19	8.02	5.71×10^{-3}	7
	VCC	7.63	7.40	1.24×10^{-2}	8
U-shaped	ACD	7.47	7.26	5.36×10^{-4}	8
	VCC	7.95	7.96	4.61×10^{-2}	5
Flappy	ACD	21.40	26.30	2.98×10^{-4}	11
	VCC	21.30	25.40	5.64×10^{-4}	13

VCC gives a higher median minimum clearance in every successful map family. This suggests that visibility-based regions often keep the rounded trajectory farther from obstacles. However, this safety margin is not free: it comes with higher decomposition cost and lower robustness on narrow passages. The relaxation and rounding objectives also vary by topology. On **flappy**, the relaxation and rounded objectives are nearly identical, indicating a tight relaxation for this corridor-like structure. On **clustering**, VCC exhibits a larger relaxation–rounding gap, showing that a compact continuous relaxation does not always preserve solution quality after discrete path selection.

4.5 Visual Evidence on Flappy Maps

The **flappy** family provides a useful qualitative check because each instance forces a winding route through alternating upper and lower openings. Figure 2 compares ACD and VCC on all five **flappy** seeds. Both methods find feasible trajectories in every seed, and their path-quality metrics are close. The main difference is computational: the median end-to-end runtime is 1.13 s for ACD and 12.90 s for VCC, mostly because VCC decomposition dominates the pipeline.

4.6 Discussion and Limitations

The main conclusion is that region count alone is not sufficient to characterize GCS difficulty. A useful decomposition must be evaluated through the induced edge count, presolved model size, relaxation and rounding behavior, decomposition time, and final trajectory quality. Under the tested one-query pipeline, ACD is the more reliable and faster default. VCC becomes compelling when the decomposition can be amortized across repeated solves or when the environment resembles **clustering**, where the visibility graph substantially reduces the GCS solve time.




Fig. 2. ACD and VCC comparisons on the five **flappy** seeds. Obstacles are black, convex regions are colored, and the GCS trajectory is plotted in blue.

The study has several limitations. Each map family contains only five seeds, and VCC has only three successful **narrows** runs. All experiments use the same workspace scale, obstacle density, obstacle size, and area-coverage mode. The reported values are extracted from diagnostic outputs; metrics not shown in those diagnostics, such as mean clearance and smoothness proxies, are not analyzed. A broader evaluation should vary workspace scale, obstacle density, VCC sampling budget, coverage target, and the number of seeds.

5 Conclusion and Future Work

Acknowledgments. We acknowledge Ho Chi Minh City University of Technology (HCMUT), VNU-HCM for supporting this study.

References

1. Deits, R., Tedrake, R.: Computing large convex regions of obstacle-free space through semidefinite programming. In: *Algorithmic Foundations of Robotics XI*, Springer Tracts in Advanced Robotics, vol. 107. Springer (2015). https://doi.org/10.1007/978-3-319-16595-0_7, https://groups.csail.mit.edu/robotics-center/public_papers/Deits14.pdf
2. Kavraki, L.E., Svestka, P., Latombe, J.C., Overmars, M.H.: Probabilistic roadmaps for path planning in high-dimensional configuration spaces. *IEEE Transactions on Robotics and Automation* **12**(4), 566–580 (1996). <https://doi.org/10.1109/70.508439>
3. LaValle, S.M.: Rapidly-exploring random trees: A new tool for path planning. Tech. Rep. TR 98-11, Iowa State University (1998)
4. Lien, J.M., Amato, N.M.: Approximate convex decomposition of polygons. *Computational Geometry: Theory and Applications* **35**(1–2), 100–123 (2006). <https://doi.org/10.1016/j.comgeo.2005.10.005>
5. Marcucci, T., Petersen, M., von Wrangell, D., Tedrake, R.: Motion planning around obstacles with convex optimization. *Science Robotics* **8**(84) (2023). <https://doi.org/10.1126/scirobotics.adf7843>, <https://arxiv.org/abs/2205.04422>
6. Marcucci, T., Umenberger, J., Parrilo, P.A., Tedrake, R.: Shortest paths in graphs of convex sets. *SIAM Journal on Optimization* **34**(1), 507–532 (2024). <https://doi.org/10.1137/22M1523790>, <https://arxiv.org/abs/2101.11565>
7. Schulman, J., Duan, Y., Ho, J., Lee, A., Awwal, I., Bradlow, H., Pan, J., Patil, S., Goldberg, K., Abbeel, P.: Motion planning with sequential convex optimization and convex collision checking. *The International Journal of Robotics Research* **33**(9), 1251–1270 (2014). <https://doi.org/10.1177/0278364914528132>
8. Werner, P., Amice, A., Marcucci, T., Rus, D., Tedrake, R.: Approximating robot configuration spaces with few convex sets using clique covers of visibility graphs. In: *2024 IEEE International Conference on Robotics and Automation (ICRA)*. pp. 10359–10365 (2024). <https://doi.org/10.1109/ICRA57147.2024.10610005>, <https://arxiv.org/abs/2310.02875>
9. Zucker, M., Ratliff, N., Dragan, A.D., Pivtoraiko, M., Klingensmith, M., Dellin, C.M., Bagnell, J.A., Srinivasa, S.S.: CHOMP: Covariant Hamiltonian optimization for motion planning. *The International Journal of Robotics Research* **32**(9–10), 1164–1193 (2013). <https://doi.org/10.1177/0278364913488805>